


DISCOVER DATABASE



Bin Zhu



WHAT'S THE DATABASE


A **database** is an organized collection of data. It is the collection of schemas, tables, queries, reports, views, and other objects. The data are typically organized to model aspects of reality in a way that supports processes requiring information

(copied from Wikipedia)



THE OBJECTS IN DATABASE

- Tables
- View
- Materialized View
- Index
- Procedure
- Sequence
- Trigger
- DB Link



TABLE

A table is a collection of related data held in a structured format within a database. It consists of columns, and rows.

Columns stores a specific data type

Row →
Or record

Emp No	Name	Age	Department	Salary
001	Alex S	26	Store	5000
002	Golith K	32	Marketing	5600
003	Rabin R	31	Marketing	5600
004	Jons	26	Security	5100



VIEW

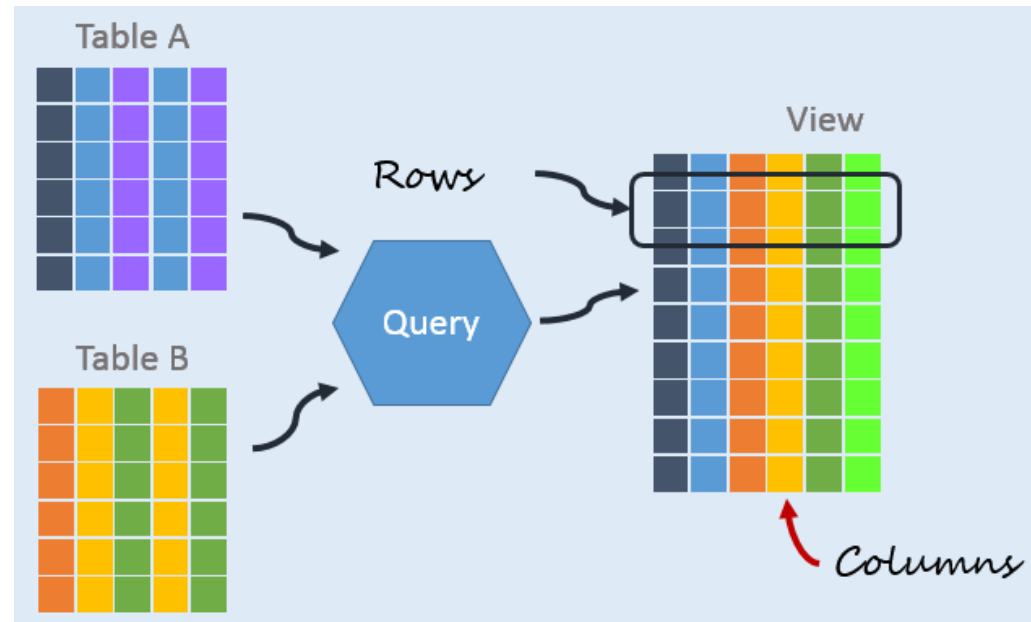
A **view** is a virtual table based on the result-set of an SQL statement.

A **view** contains rows and columns, just like a real table.

The fields in a **view** are fields from one or more real tables in the **database**.

Sample:

```
CREATE VIEW view_dept_201  
AS  
SELECT emp_id, name, hire_date  
FROM gdb.employees  
WHERE department = 201;
```




MATERIALIZED VIEW

A **materialized view** is a database object that contains the results of a query

Sample:

Create materialized view myView as select * from myTab;

Refresh materialized view myView;



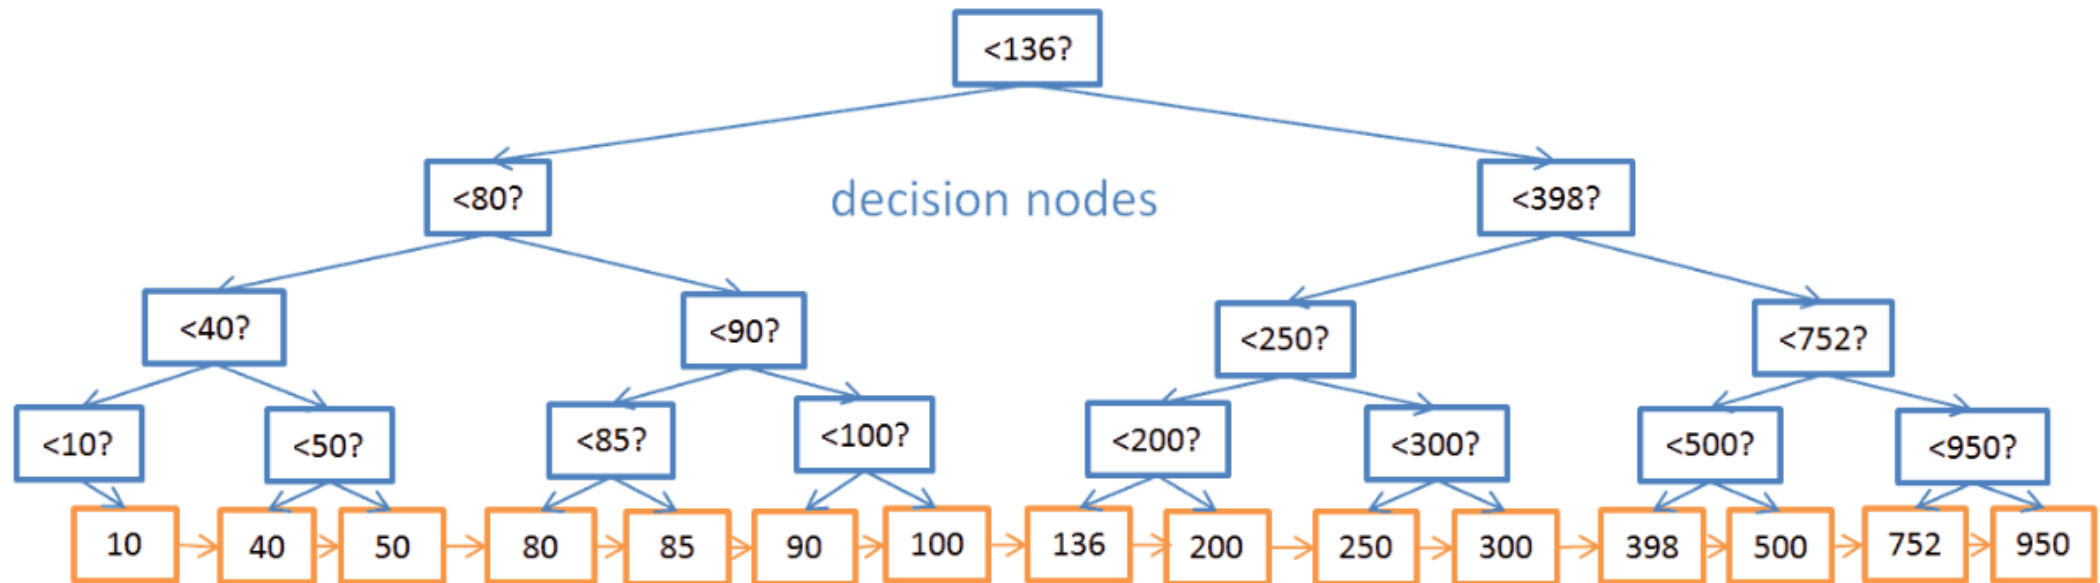
INDEX

- Indexes are **Database Object** that you can create to improve the performance of query.
- Indexes can be created **explicitly** or **automatically**.
- It is used to **retrieve data faster**.
- You can create indexes on one or more columns of a table.
- You can create more than one index on one table.
- User can't see indexes.

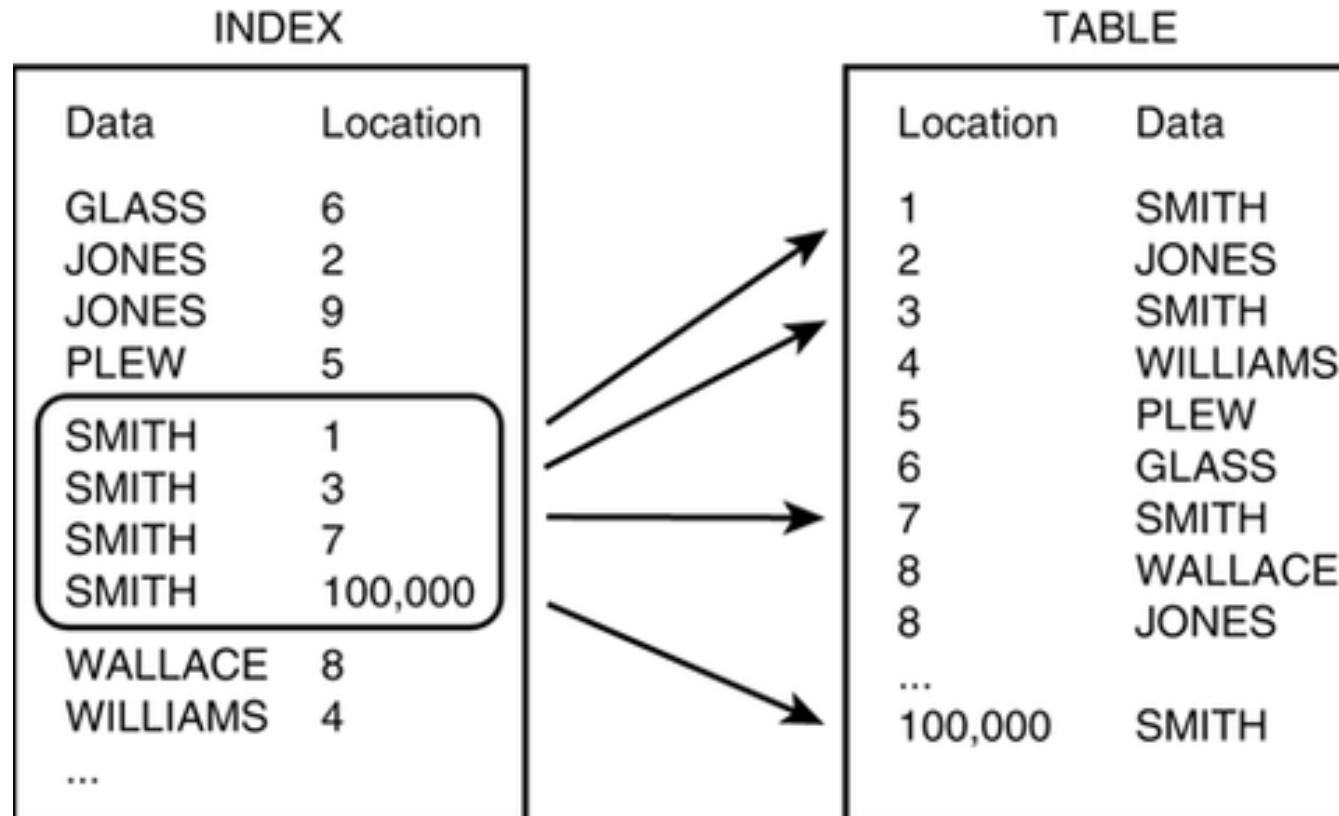


INDEX – HOW IT WORKS

B+ Tree / Database index

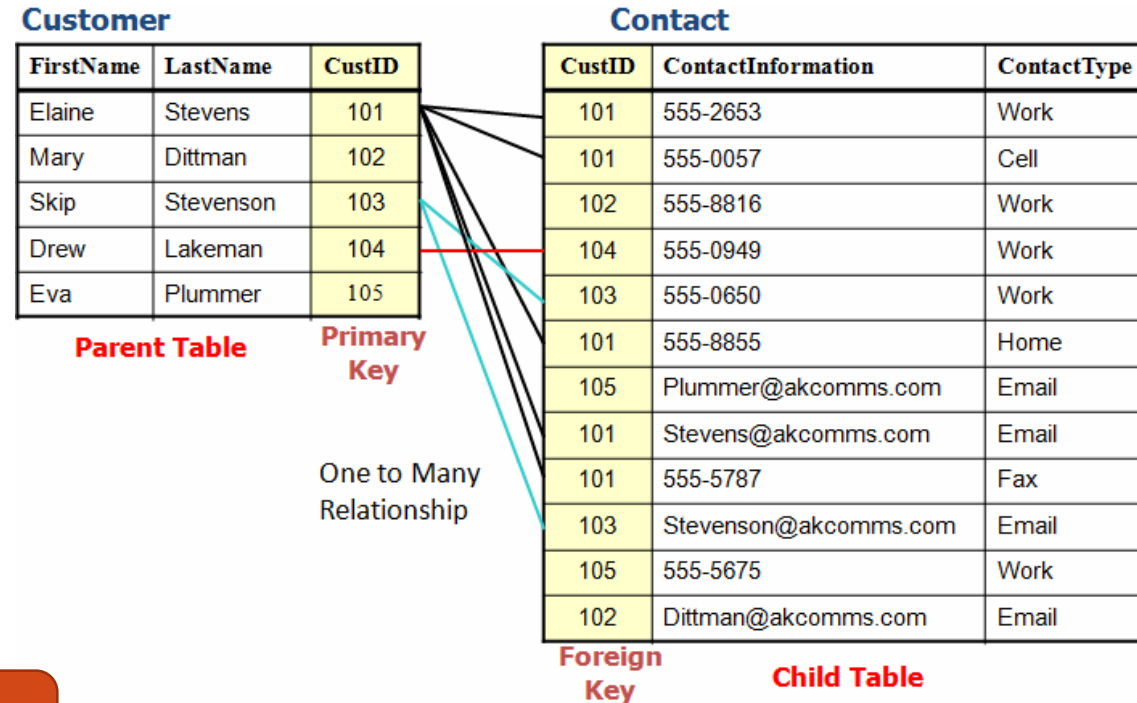


INDEX – HOW IT WORKS



INDEX - TYPE

- Primary Key
- Unique Index
- Foreign key
- Index



Primary Key



Unique Index



PROCEDURE

A **stored procedure** is a set of Structured Query Language (SQL) statements with an assigned name that's **stored** in the **database** in compiled form so that it can be shared by a number of programs.

Sample


```
ALTER PROCEDURE <ProcedureName>
    -- Add the parameters for the stored procedure here
    <@Parameter1> <Datatype_For_Parameter1> = <Default_Value>,
    <@Parameter1> <Datatype_For_Parameter1> = <Default_Value>
AS
BEGIN
    -- Insert statements for procedure here
    SELECT * FROM TableName
END
```



SEQUENCE

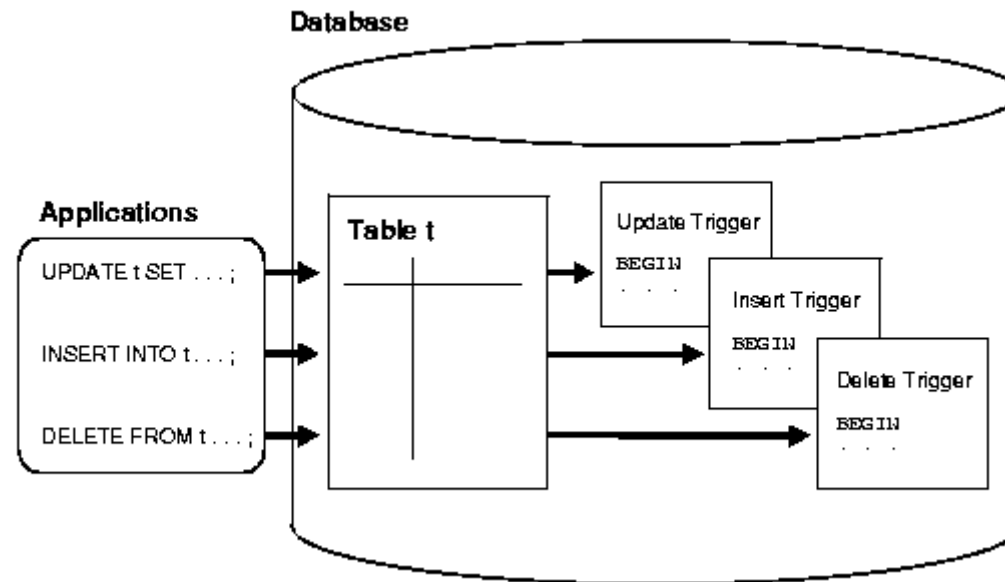
A sequence:

- Can automatically generate unique numbers
- Is a shareable object
- Can be used to create a primary key value
- Replaces application code
- Speeds up the efficiency of accessing sequence values when cached in memory




TRIGGER

A **trigger** is a special kind of stored procedure that automatically executes when an event occurs in the **database** server. DML **triggers** execute when a user tries to modify data through a data manipulation language (DML) event. DML events are INSERT, UPDATE, or DELETE statements on a table or view



DB LINK

Dblink is a module which supports connections to other databases from within a database session



DATABASE TRANSACTION

- ACID
- Transaction
- Distributed Transaction
- Two Phase Commit
- Real Commit Diagram



DATABASE BASIC PROPERTIES

A

Atomicity requires that each transaction be "all or nothing": if one part of the transaction fails, then the entire transaction fails, and the database state is left unchanged.

C

The consistency property ensures that any transaction will bring the database from one valid state to another.

I

The isolation property ensures that the concurrent execution of transactions results in a system state that would be obtained if transactions were executed serially, i.e., one after the other.

D

The durability property ensures that once a transaction has been committed, it will remain so, even in the event of power loss, crashes, or errors.



TRANSACTION

A transaction is a unit of work that you want to treat as "a whole". It has to either happen in full, or not at all.



Commit

OR

Rollback


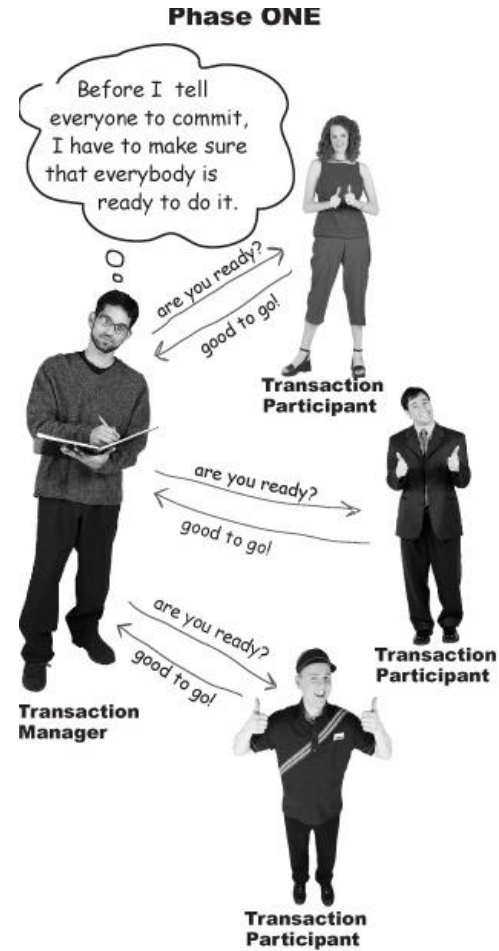


DISTRIBUTED TRANSACTION


A **distributed transaction** is a database **transaction** in which two or more network hosts are involved



TWO PHASE COMMIT

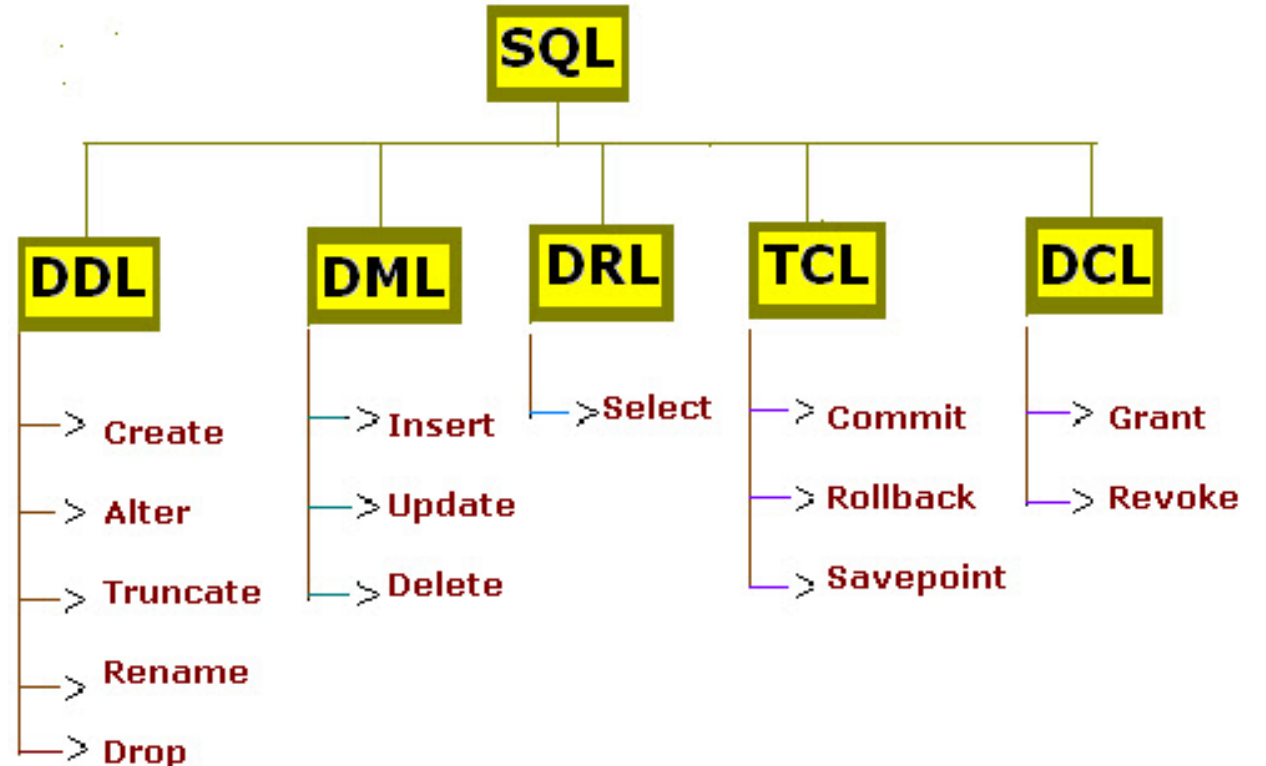


DATABASE REAL COMMIT DIAGRAM




DATABASE SQL OPERATION

- DDL:
Data Definition Language
- DML:
Data Manipulation Language
- DRL:
Data Retrieval Language
- TCL:
Transaction Control Language
- DCL:
Data Control Language



DATABASE DEADLOCK



A **deadlock** is a situation in which two or more transactions are waiting for one another to give up locks.



- 1, **Select for Update**
- 2, Insert
- 3, Update
- 4, Delete



DATABASE PERFORMANCE




IO TUNING PERFORMANCE


- Tablespace

A **tablespace** is a storage location where the actual data underlying database objects can be kept. It provides a layer of abstraction between physical and logical data, and serves to allocate storage for all DBMS managed segments.

- Fragmentation



TABLESPACE



FRAGMENTATION

- Table Fragmentation
- Index Fragmentation
- How to Perform the fragmentation
 - Vacuum Database
 - Rebuild Index
 - Rebuild Table




TABLE PARTITION

A **partition** is a division of a logical database or its constituent elements into distinct independent parts. Database partitioning is normally done for manageability, performance or availability reasons, as for load balancing.

Partitioning criteria

- Range partitioning
- List partitioning
- Hash partitioning
- Composite partitioning




RANGE, LIST, HASH PARTITIONING

List
Partitioning



Range
Partitioning




Hash
Partitioning




COMPOSITE PARTITIONING


Composite Partitioning
Range-Hash



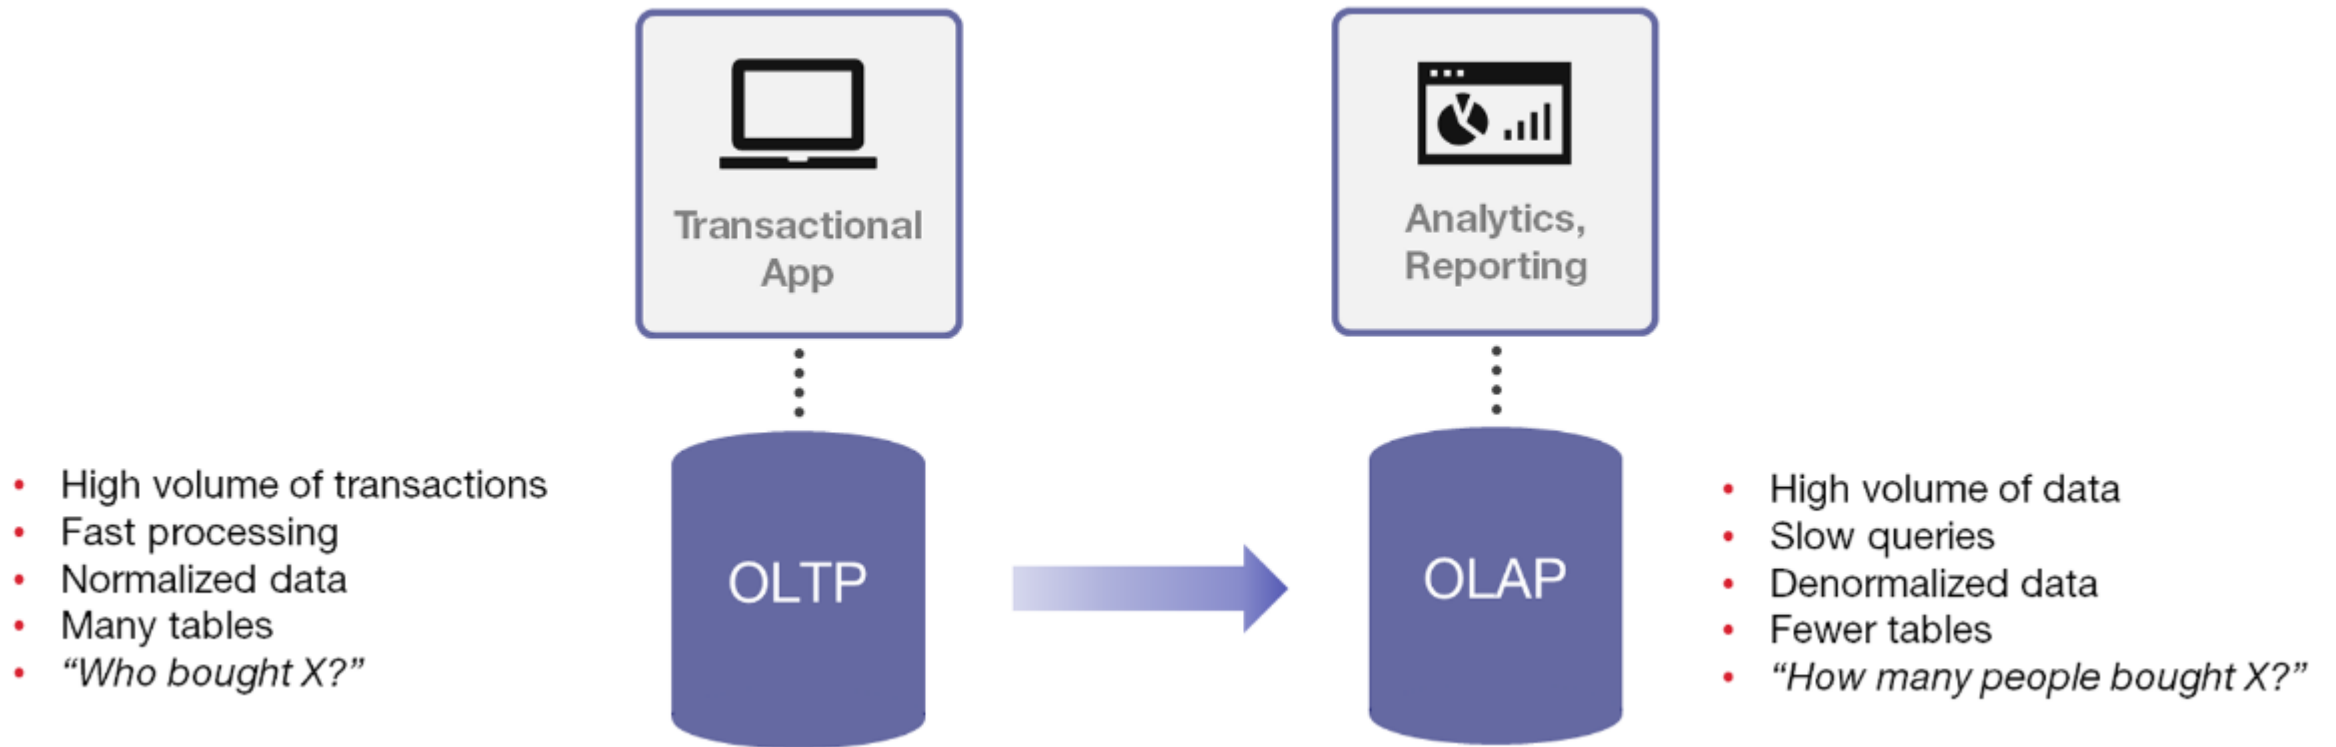
Composite Partitioning
Range - List



PARTITION INDEX

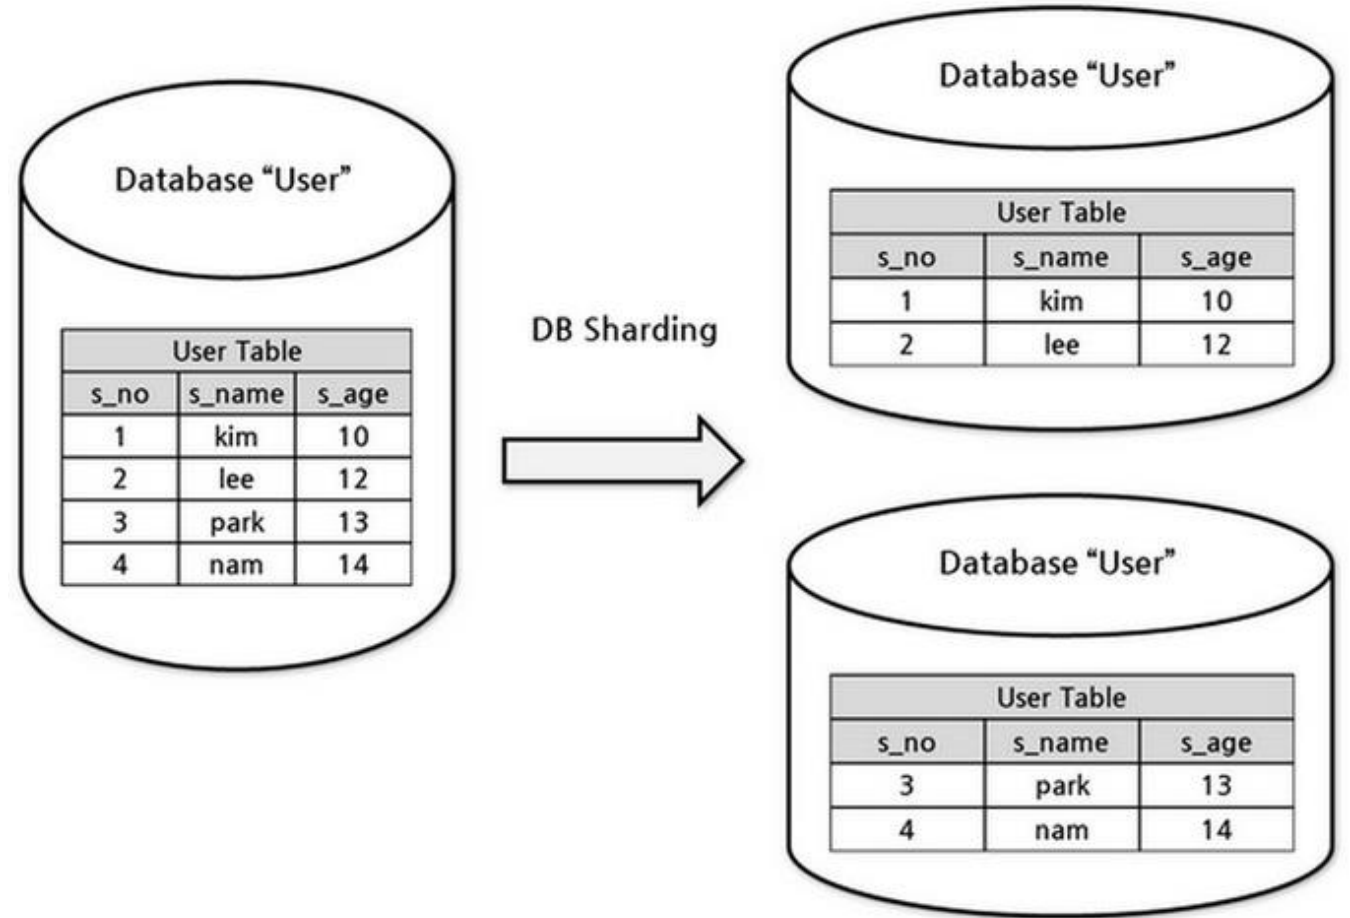


OLTP VS OLAP




DATABASE SHARDING

Database Sharding is a highly scalable approach for improving the throughput and overall performance of high-transaction, large database-centric business applications



DATABASE SHARDING

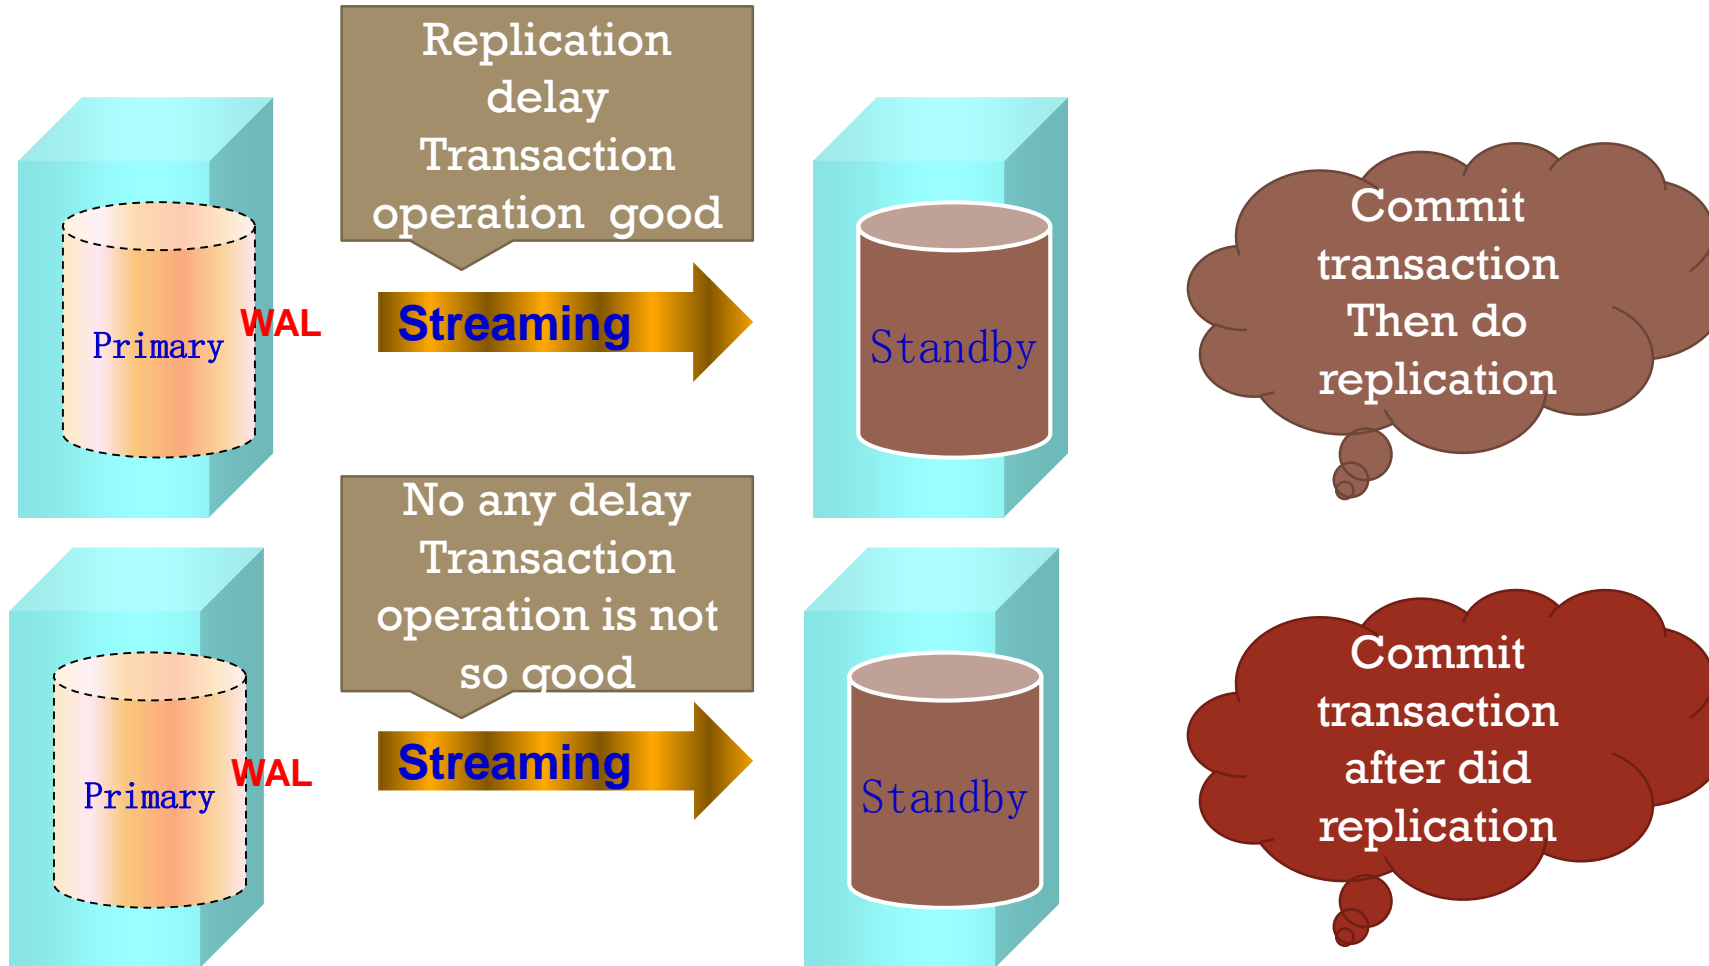


DATABASE REPLICATION

Database replication is the frequent electronic copying data from a database in one computer or server to a database in another so that all users share the same level of information.



ASYNC REPLICATION VS SYNC



DATABASE CLUSTER

Clustering offers two major advantages, especially in high-volume database environments:

- **Fault tolerance:** Because there is more than one server or instance for users to connect to, clustering offers an alternative, in the event of individual server failure.
- **Load balancing:** The clustering feature is usually set up to allow users to be automatically allocated to the server with the least load.




DATABASE CLUSTER TYPE

- Shared Data Array
- Master – Slave(s)
- Master - Master





SHARED DATA ARRAY



- ❖ Positive
Shared data disk array solution can avoid take long time to synchronized the data between the different database. When the postgres instance1 failed, you just need to assign all the clients to connect to instance2 and all the data will not lost.
- ❖ Negative
Because all the data is stored in the same disks, if the disks are failed, the data will lost.

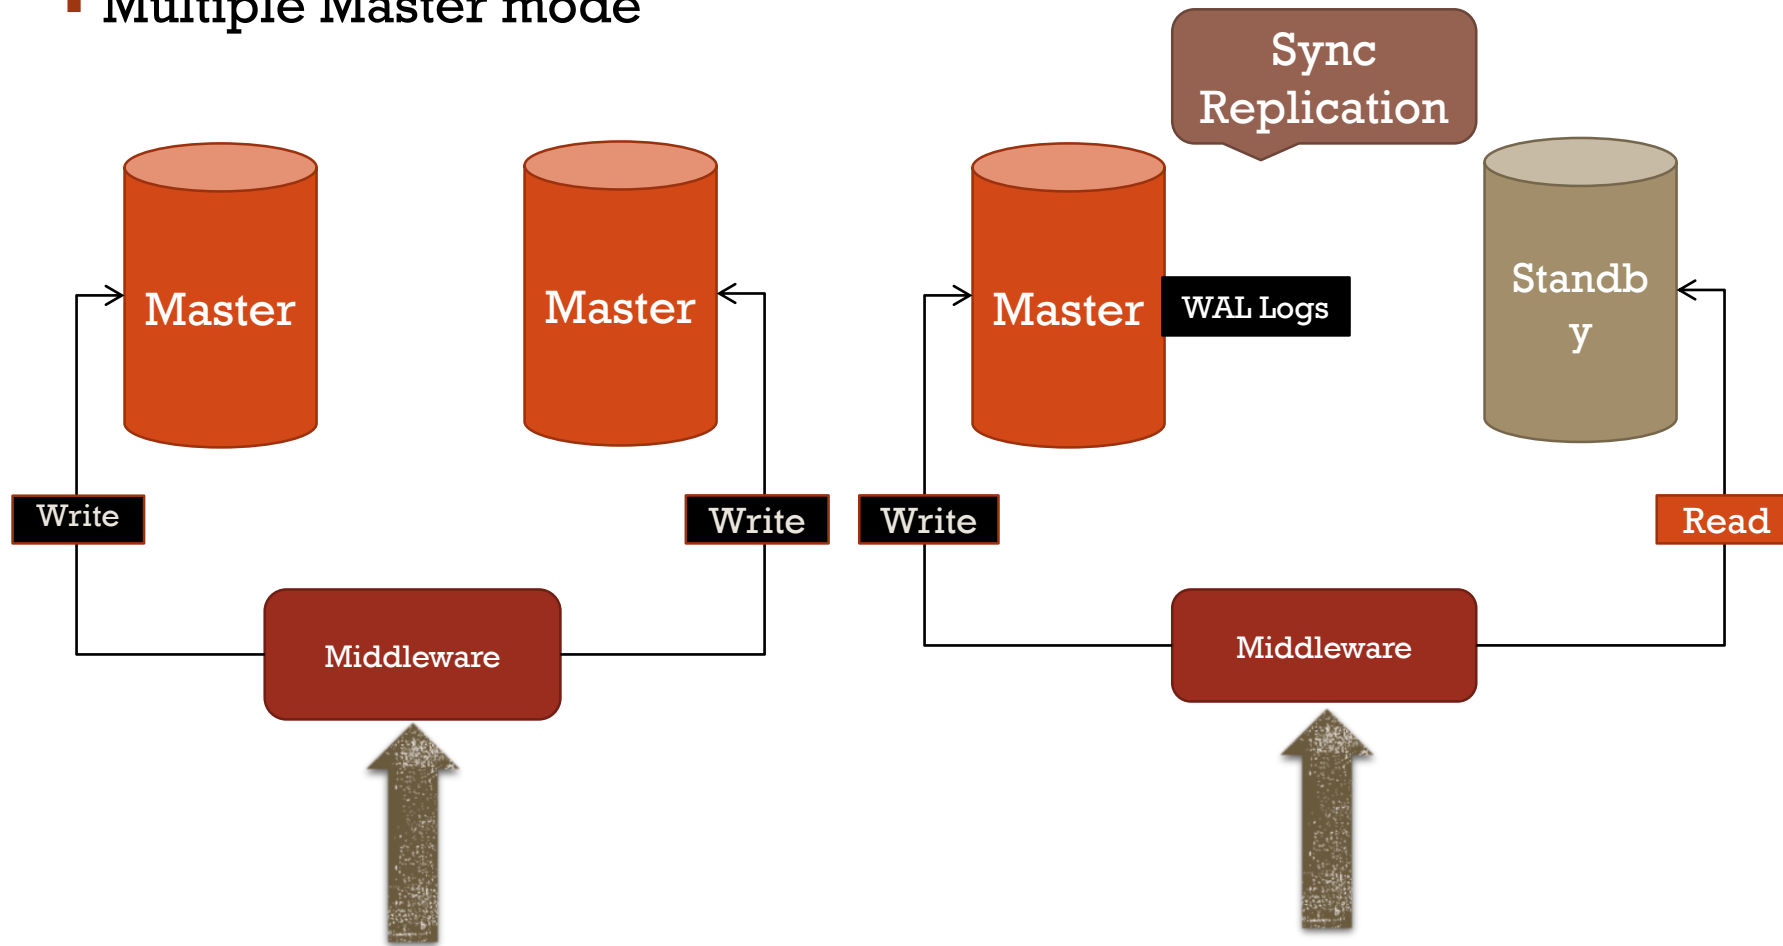


MASTER — SLAVE(S)



MASTER – MASTER

- Multiple Master mode



CLOUD DATABASE




Google
Cloud Platform



Amazon RDS Database Engines



CLOUD DATABASE VIDEO (COPIED FROM AWS)



RDBM VS NOSQL

No-SQL

Features

- Non-Relational DB Model
- Horizontal Scaling
- CAP Theorem Amenability
- Schema Creation Possible at runtime



Anyone Using NoSQL Model?

- Google, Amazon



Pros of NoSQL

- Ability to store complex data type
- No significant changes in code when data structure is altered.



Cons of NoSQL

- Still at early stage of evolution
- NO Standardized model available



SQL

Features

- Relational DB Model
- Vertical Scaling
- ACID Amenability
- Schema Creation and locking before data entry



Anyone Using SQL Model?

- Most of the firms are currently using SQL Model



Pros of SQL

- Ability to support Transactions
- Ability to support powerful aggregation of data



Cons of SQL

- Not able to handle big data
- Management is vendor support dependant



MPowersoft_Comm©



VIDEO (DYNAMO) (COPIED FROM AWS)



NEW SQL

NewSQL is a class of modern relational database management systems that seek to provide the same scalable performance of NoSQL systems for online transaction processing (OLTP) read-write workloads while still maintaining the ACID guarantees of a traditional database system

-- Wikipedia



VIDEO (GOOGLE SPANNER) (COPIED FROM GOOGLE)



CATALOGUE

- ❖ What's the database
- ❖ The objects in database
 - Tables
 - View
 - Materialized View
 - Index
 - Procedure
 - Sequence
 - Trigger
 - DB Link
- ❖ Database Transaction
 - ACID
 - Transaction
 - Distributed Transaction
 - Two phase commit
 - Database Real Commit Diagram
- Database operation
 - DDL
 - DML
 - DRL
 - TCL
 - DCL
- ❖ Database deadlock
- ❖ Database performance
 - Tuning performance
 - Tablespace
 - Fragmentation
- ❖ Table partition
 - Range
 - List
 - Hash
 - Composite
- Partition index
 - Global
 - Local
- ❖ OLTP VS OLAP
- ❖ DB sharding
- ❖ DB replication
 - Async
 - Sync
- ❖ DB Cluster
 - Shared Disk Array
 - Master – Slave(s)
 - Multiple Master
- ❖ Cloud Database
- ❖ RDBM VS NoSQL
- ❖ NewSQL



